

How to Size Trades in a Leveraged Macro Portfolio

Sizing is often the difference between a good idea and a viable trade.

Published: 2026-04-28 | Tags: Notes • Portfolio Construction • Risk Management • Macro

A personal educational framework note on turning trade ideas into robust risk allocations under leverage, volatility, and portfolio constraints.

Trade sizing is usually treated as a secondary question. The thesis comes first. The entry level comes first. The narrative comes first.

But in a leveraged macro portfolio, sizing is the point at which a trade is either made investable or rendered dangerous. A good idea sized badly can become unmanageable under stress. A more modest idea sized well can survive long enough to become useful. The difference shows up in drawdowns, rebalancing pressure, liquidity constraints, and the ability to maintain objectivity when the market moves against the position.

This note is not a prescription for position sizing. It is a framework for thinking about how trade ideas become risk allocations under uncertainty.

A robust sizing framework should answer a single clarifying question:

How much risk can this trade carry, given the setup, the expected path, the historical distribution, the scenario range, and the rest of the book?

1. Assess the Path, Not Just the Destination

The first sizing decision is shaped by the expected path of the trade.

Most discussions focus on expected return: where the asset should go if the thesis is correct. But a portfolio manager also has to ask how it might get there. A trade with asymmetric upside but an unstable entry path may require a smaller initial allocation than a trade with lower upside but cleaner path dependency.

In leveraged portfolios, mark-to-market losses can force reductions before the underlying thesis has had time to resolve. The relevant question is not only whether the trade is attractive at the target horizon. It is whether the position can survive the path required to reach that horizon.

A sizing process should therefore start with path risk: the likely drawdown profile, the expected volatility around entry, the liquidity available under stress, and the extent to which adverse price action would invalidate the thesis or merely test it.

2. Use the Chart as a Map, Not a Forecast

Technical analysis is often overstated, but levels matter because market participants risk capital around them.

Support, resistance, retracements, prior highs and lows, moving-average zones, and crowded stop areas all influence short-term behaviour. They matter not because they predict the future, but because they reveal where positioning, institutional memory, liquidity and path dependency may


concentrate.

A chart can help identify:

- where liquidity might deteriorate
- where stop-loss behaviour may accelerate
- where a better entry level might justify adding risk
- where the original trade location starts to look poor
- where the market is likely to test conviction

The chart should not decide the thesis. It should help shape the first clip, the stop distance, and the amount of capital held in reserve.

For sizing, this distinction is important. A chart is not a forecasting engine. It is a map of where the trade may become difficult to hold.



US 10-year yields remained in a well-defined downward channel for decades before breaking higher in 2022; since then, yields have spent several years around the 23.6% retracement zone.

The point is not that a retracement level has predictive power by itself. The point is that long-lived technical structures can shape how investors frame risk. A level that has acted as a reference point for years may influence entry timing, stop placement, hedge construction and the amount of risk a manager is willing to add after a move has already occurred.

For trade sizing, the chart is useful because it defines the terrain.

It does not tell the manager what will happen next. It helps describe where the position may become uncomfortable, where market attention may cluster, and where the trade may need to be reassessed.

3. Fading Stress: Act Early, Size Modestly, Leave Ammunition

Fading an extreme market stress can be structurally attractive but psychologically difficult. Volatility is high, price action is disorderly, and the mark-to-market risk is immediate.

The sizing problem is survival.

When volatility is elevated, the first allocation should usually be smaller than normal. The potential upside may already be large relative to ordinary daily variance, but the risk of further disorder is also high. The worst outcome is to identify the opportunity correctly, size it too aggressively, and be forced to reduce exposure just as the dislocation reaches its most attractive point.

A useful principle is:

act early, size modestly, and leave ammunition.

Incremental sizing matters. Later additions must be large enough to improve the average entry level meaningfully, not just cosmetically. If the initial position consumes too much risk budget, the manager loses the ability to respond when the market creates a better level.

A stress fade should be sized around survivability first and upside second.

4. Chasing Momentum: Size for the Retracement

When a move is already in motion, the sizing problem changes.

The issue is no longer catching a falling knife. It is overpaying for participation.

Momentum trades often look cleanest after a large part of the move has already occurred. The thesis may be improving, the price action may be confirming, and the market narrative may be broadening. But the entry quality is usually deteriorating at the same time.

A momentum entry should therefore be sized around the expected test. Even good trends retrace. A position that cannot withstand a normal pullback is not really sized for the trade; it is sized for immediate continuation.

The goal is to own enough risk to participate if the move extends, but not so much that a healthy retracement forces an exit.

This is especially important in macro markets, where trend, carry, valuation and policy narrative can point in the same direction for a period before reversing sharply. Momentum can be a valid reason to own risk. It is rarely a valid reason to ignore entry quality.

5. Joining a Sustained Trend: Respect the Crowding

Durable trends often extend far beyond what valuation metrics suggest is reasonable. A currency can stay expensive, a curve can keep steepening, a spread can keep tightening, and a rates market can continue repricing long after the obvious fundamental argument has been made.

The longer the trend, however, the more the sizing question shifts from thesis quality to crowding risk.

Crowded trades behave poorly when volatility rises. Sharp counter-moves may not reflect a change in the macro thesis, but they can still trigger stops, reduce liquidity, and force de-risking. In those environments, the price action can become more about positioning than information.

This argues for partial risk sizing, wider tolerance for noise, and a willingness to trim exposure into disorderly extensions. The aim is not to avoid trends. Trends are often where the largest macro returns are made. The aim is to recognise that late trend entry has a different risk profile from early trend formation.

A sustained trend should be sized with respect for the exit, not just enthusiasm for the thesis.

6. Calibrate with VaR, but Do Not Worship It

Risk measures such as Value at Risk are regime-dependent and fragile in the tails. They are backward-looking, sensitive to the lookback window, and least reliable when the future differs most from the recent past.

They are still useful.

VaR provides a numerical yardstick. It helps compare one trade with another, converts different assets into a common risk language, and highlights when a position is large merely because trailing volatility appears low.

A good sizing process uses VaR to ask:

- how large is this trade relative to recent volatility?
- how does the risk compare with other positions in the portfolio?
- is the position large because conviction is high, or because measured volatility is artificially suppressed?
- what happens if the volatility regime changes?
- does the trade still make sense if the risk estimate doubles?

VaR is a ruler, not a religion.

It is most useful when combined with stress scenarios, liquidity assumptions, and a clear view of the current volatility regime.

7. Analyze History for Distributions, Not Just Valuations

Historical analysis is often used to support valuation claims: this spread is wide, this curve is steep, this asset is cheap, this relationship has mean-reverted before.

That can be useful, but it is not enough.

Valuation precedents are broken with surprising frequency. The more important use of history is distributional: how the asset moves, how often it gaps, how fat the tails are, and how it behaves when broader risk conditions change.

Instead of asking only whether a market has traded at a more extreme valuation before, the sizing process should ask:

- what is a normal weekly move?
- what is a two-standard-deviation move?
- how quickly do losses tend to accumulate?
- are the tails symmetric or skewed?
- how does the asset behave when liquidity deteriorates?
- does the trade diversify other portfolio risks, or does it lose money at the same time as everything else?

This is where apparently similar trades can differ materially.

Two trades may have the same expected return and similar headline volatility, but one may lose money gradually while the other gaps through the exit level. One may diversify the book in a risk-off shock while the other quietly amplifies the same factor exposure already present elsewhere.

Sizing should reflect that difference.


8. Parameterize the Risk

Historical analysis shows what has happened. Scenario analysis asks what could happen.

A useful way to build scenarios is to decompose a trade into its underlying valuation drivers. Instead of treating the position as a single price or spread, the portfolio manager can separate the building blocks that explain why it makes or loses money.

For example, a five-year Italian government bond exposure might be decomposed into:

1. Terminal rates — proxied by forward swaps, such as 1y1y.
2. Term premium — proxied by the curve slope, such as 2s5s.
3. Sovereign spread — proxied by the spread versus OIS or core government comparables.



Example of a non-linear relationship between curve structure and credit spread regimes: the 2s5s10s fly can behave with a convex profile across different spread environments.

The example illustrates why decomposition matters. Some trades do not behave linearly across regimes. A curve fly may appear to have one risk profile when credit spreads are stable and a very different profile when sovereign or credit risk begins to dominate the rates market.

In this type of structure, the question is not simply whether the fly is cheap or rich on a standalone basis. The more useful question is how the fly behaves as the macro regime changes. If the relationship between curve structure and spread levels is non-linear, then the sizing process has to account for that convexity.

This decomposition makes the risk more legible.

A loss can then be attributed to a rates shock, a curve shock, a spread shock, or a combination of all three. It also allows the manager to build synthetic scenarios that have not appeared cleanly in recent history.

For example:

- terminal rates rise while sovereign spreads are stable
- spreads widen while the rates market rallies
- the curve steepens while credit risk reprices

- all components move adversely at once
- the hedge works on the rates component but fails on the spread component

This type of scenario work is more useful than asking whether a trade is cheap in isolation. It clarifies what the position is really short, what it is really long, and which proxy hedges are likely to work under stress.

9. Establish Explicit P&L Bounds

Once the scenarios are defined, the trade needs explicit P&L bounds before execution.

This does not mean pretending that losses can be predicted precisely. It means defining the range of outcomes that the position is expected to tolerate.

A practical sizing framework should distinguish between:

- a normal adverse move
- a severe but plausible adverse move
- a favourable base case
- a stress outcome
- the condition that conceptually invalidates the trade

The difference between a painful mark-to-market move and a broken thesis should be clear before the trade is entered. Without that distinction, every loss becomes ambiguous. The manager either exits too early because the position is uncomfortable, or averages down too easily because the original thesis has not been formally invalidated.

Good sizing reduces that ambiguity.

A normal adverse move should not force the trade out of the portfolio. A severe adverse move should be survivable if the thesis remains intact. But the framework should not be so loose that every loss is treated as an opportunity to add.

The purpose of P&L bounds is not to eliminate judgement. It is to make judgement less emotional.

10. Assess Portfolio Fit and Marginal Risk

Trade-level sizing is incomplete without portfolio context.

A position that looks sensible on a standalone basis may be too large once its hidden factor exposures are included. A rates trade may be another expression of the same dollar risk already present in FX. A sovereign-spread trade may increase the same risk-off exposure embedded in credit. A curve trade may look diversified until it is stress-tested against the portfolio's broader policy-rate exposure.

The final allocation should therefore pass three checks.

First, marginal risk contribution. Does the trade diversify the portfolio, or does it amplify an existing exposure?

Second, opportunity cost. Does the conviction and payoff quality justify using scarce capital, liquidity, risk limits and attention?

Third, regime context. Is the current market environment likely to reward the edge being expressed, or is the trade fighting the dominant volatility, liquidity or policy regime?

Only after these checks does a trade idea become a risk allocation.

The objective is not to size every position perfectly. That is impossible. The objective is to avoid making the most common mistake: sizing a trade as if it exists in isolation.

Conclusion

Trade sizing is where research meets portfolio reality.

A thesis can be thoughtful, valuation can be attractive, and the macro logic can be sound. None of that is enough if the position is too large for the path it has to survive, too fragile for the liquidity available, or too correlated with the rest of the book.

A useful sizing framework therefore starts with the expected path, uses charts as maps of market behaviour, calibrates risk with volatility and distributional history, tests scenarios through underlying valuation drivers, and finally asks whether the position improves the portfolio as a whole.

The central question is not simply:

How much can this trade make?

It is:

How much risk should this idea be allowed to carry?

That is the point where a trade becomes a portfolio decision.

Disclaimer

This article is a personal framework note for general informational and educational purposes only. It is not investment advice, not a recommendation, and not a solicitation to buy or sell any security, financial instrument, or strategy. Any examples are simplified and illustrative. They may omit transaction costs, financing costs, liquidity constraints, taxes, market impact, and other practical considerations. The views expressed are personal and are not affiliated with, endorsed by, or representative of any employer or institution.